

The seal of Georgetown University is a circular emblem. It features an eagle with its wings spread, perched on a shield with vertical stripes. Above the eagle's head is a lyre. The eagle is surrounded by a wreath of leaves. The circular border of the seal contains the Latin text "MACI IN MARYLANDIA" at the top and "GEORGIOPOLITANUM" at the bottom, with stars interspersed. The motto "UTRAQUE UNUM" is written on a ribbon across the eagle's chest.

Predictive Analytics: Air Pollution

Georgetown University Certificate in
Data Science, Spring 2020

GEORGETOWN
UNIVERSITY

Team Members

1. *James Hinshaw*

2. *Ayesha Baig*

3. *Adrienne White*

4. *Jeremy Lykken*

5. *Julien Collear*

Agenda

1. Introduction

2. Methodology, Data Pipeline, Ingestion and Wrangling

3. Exploratory Data Analysis

4. Feature Engineering, Model Selection and Machine Learning


6. Results, Conclusions, Next Steps and Closing Remarks

**20
Minutes**

Introduction



Background



Motivation

9



4.2 million

150 million

Motivation

- The ability to locate, measure, and predict PM2.5 levels is vital for protecting vulnerable population within local communities
- The current state of PM2.5 monitoring
- Government Sensors
 - limited access to data
 - small network size
 - expensive
 - accurate but inconsistent monitoring



Application/Hypothesis

- The goal of this study was to develop a local, high-resolution air pollution model which can both predict changes in $PM_{2.5}$ levels and provide residents in the District of Columbia a tool for making healthier decisions about where and when they spend their time outside.
- Hypothesis
 - The accuracy of air quality predictions will be influenced by
 - spatial and temporal gaps
 - location features
 - meteorological inputs
- How are we going to do this?
 - public and private sensors
 - high resolution measurements
 - consider the impact of location, location features, and weather inputs on pm_{25} concentrations



Methodology



The input values/columns

```
cols = ['x', 'y', 'population', 'dist-mroads',  
        'dist-setl', 'dist-coast', 'dist-forest', 'slope', 'elevation',  
        'dayofweek', 'sin_day', 'cos_day', 'sin_year', 'cos_year', 'TEMP',  
        'DEW', 'SKY', 'VIS', 'ATM', 'Wind-Rate', 'sin_wind', 'cos_wind']  
v = df['pm25']
```

Classification Model Methodology

```
In [ ]: from sklearn.tree import DecisionTreeClassifier
parameters = {'criterion':["gini","entropy"],
              "max_depth":[2,3,4],
              "min_samples_leaf":[5,8,10]}
dt = DecisionTreeClassifier(random_state=0)
clf = GridSearchCV(dt, parameters)
clf.fit(X_train, y_train)
```

Regression Model Methodology

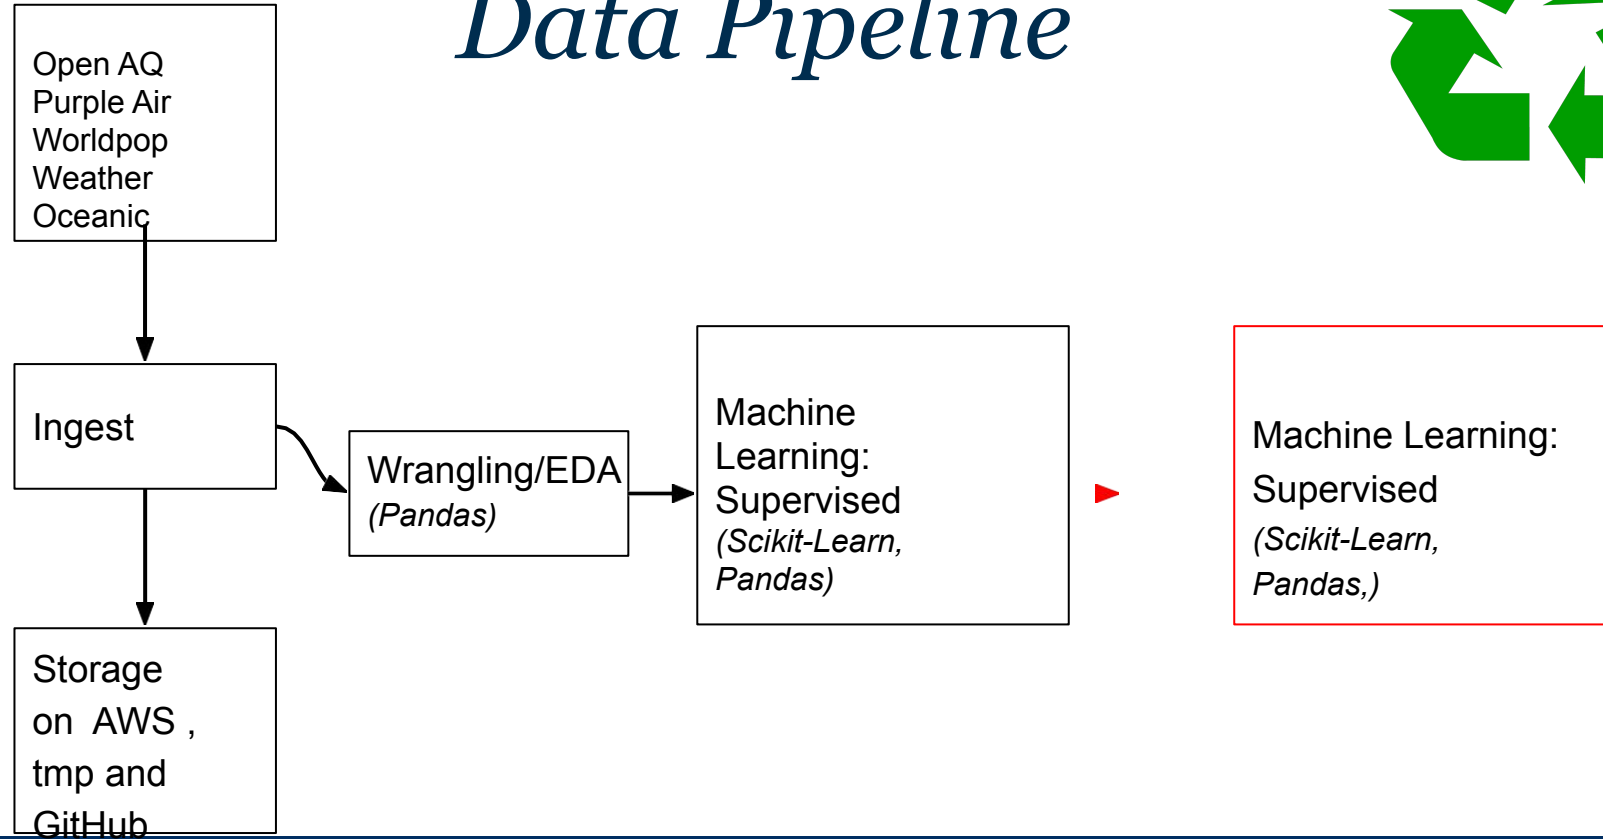
```
model_dict = {
    "DecisionTreeRegressor" : {"model":DecisionTreeRegressor,
                               "parameter":{
                                   "criterion": ["mse", "friedman_mse"],
                                   "splitter" : ["best", "random"],
                                   "min_samples_leaf" : [1,2,3]
                               },
                               'par':{
                                   "criterion": ["mse"],
                                   "splitter" : ["best", "random"],
                               }},
    "MLPRegressor" : {"model":MLPRegressor,
                      "parameter":{
                          "hidden_layer_sizes": [100,200],
                          "activation" : ['identity', 'logistic', 'tanh', 'relu'],
                          "alpha" : [0.0001,0.001]
                      },
                      'par':{
                          "hidden_layer_sizes": [100],
                          "activation" : ['identity'],
                      }},
    "KNeighborsRegressor" : {"model":KNeighborsRegressor,
                             "parameter":{
                                 "n_neighbors": [5,6],
                                 "algorithm" : ['auto', 'ball_tree']
                             },
                             'par':{
                                 "n_neighbors": [5],
                                 "algorithm" : ['auto']
                             }
    }
}

#list of models to be used
models = ['DecisionTreeRegressor', 'MLPRegressor', 'KNeighborsRegressor']
```

Data Pipeline



Data Pipeline



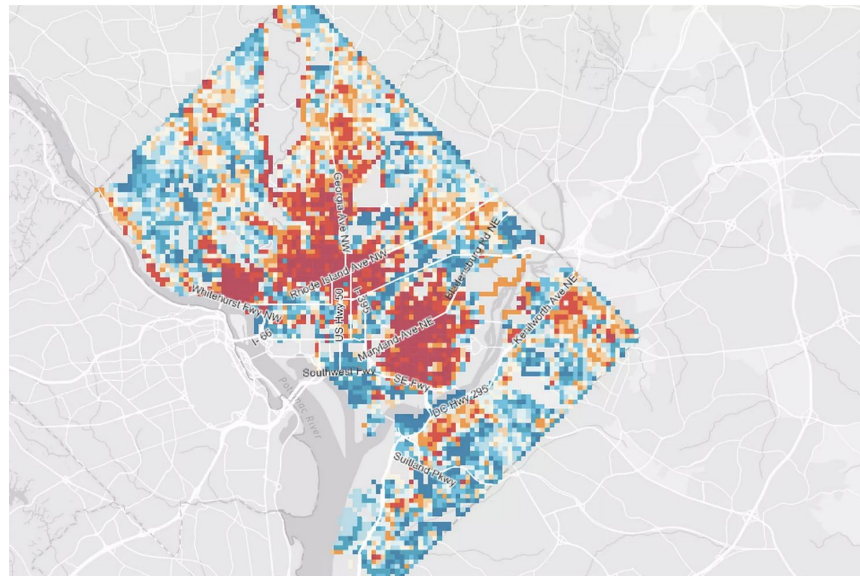
Ingestion and Wrangling




Ingestion & Wrangling



- Relied on 4 main data sources:
 - Open AQ
 - PurpleAir
 - WorldPop
 - Weather Data
- Ground sensors based on geographic locations.
 - In the DC metro area however it was expanded to beyond to get a clearer picture of the
- Switched to
 - CSV based as opposed to all the different file sources we had.






Pandas,

pandas

$$y_i = \beta^T x_i + \mu_i + \epsilon_i$$



plotly

Scikit-learn, Numpy



NumPy

Matplotlib

matplotlib

Wrangling



Weather / meteo data extract

weather data extraction made manually and storded locally

Data wrangling done here below

```
In [13]: df = pd.read_csv("bigtable.csv")
df['time'] = pd.to_datetime(df["datetime"],format="%Y-%m-%d %H:%M:%S %Z")
df['date'] = pd.Series(df.time).dt.strftime("%Y-%m-%d")
df['hour'] = pd.Series(df.time).dt.strftime("%H")

print("nb ground station measurement : {}".format(len(df)))

#2019-01-01T02:40:00
df_meteo = pd.read_csv("weather/weather.csv")
df_meteo['time'] = pd.to_datetime(df_meteo["DATE"],format="%Y-%m-%dT%H:%M:%S")
df_meteo['date'] = pd.Series(df_meteo.time).dt.strftime("%Y-%m-%d")
df_meteo['hour'] = pd.Series(df_meteo.time).dt.strftime("%H")
df_meteo['datehour'] = pd.Series(df_meteo.time).dt.strftime("%Y-%m-%d-%H")

print("nb meteo : {}".format(len(df_meteo)))

df_meteo.head()
```

nb ground station measurement : 152624

nb meteo : 14449




EDA
Exploratory
Data
Analysis

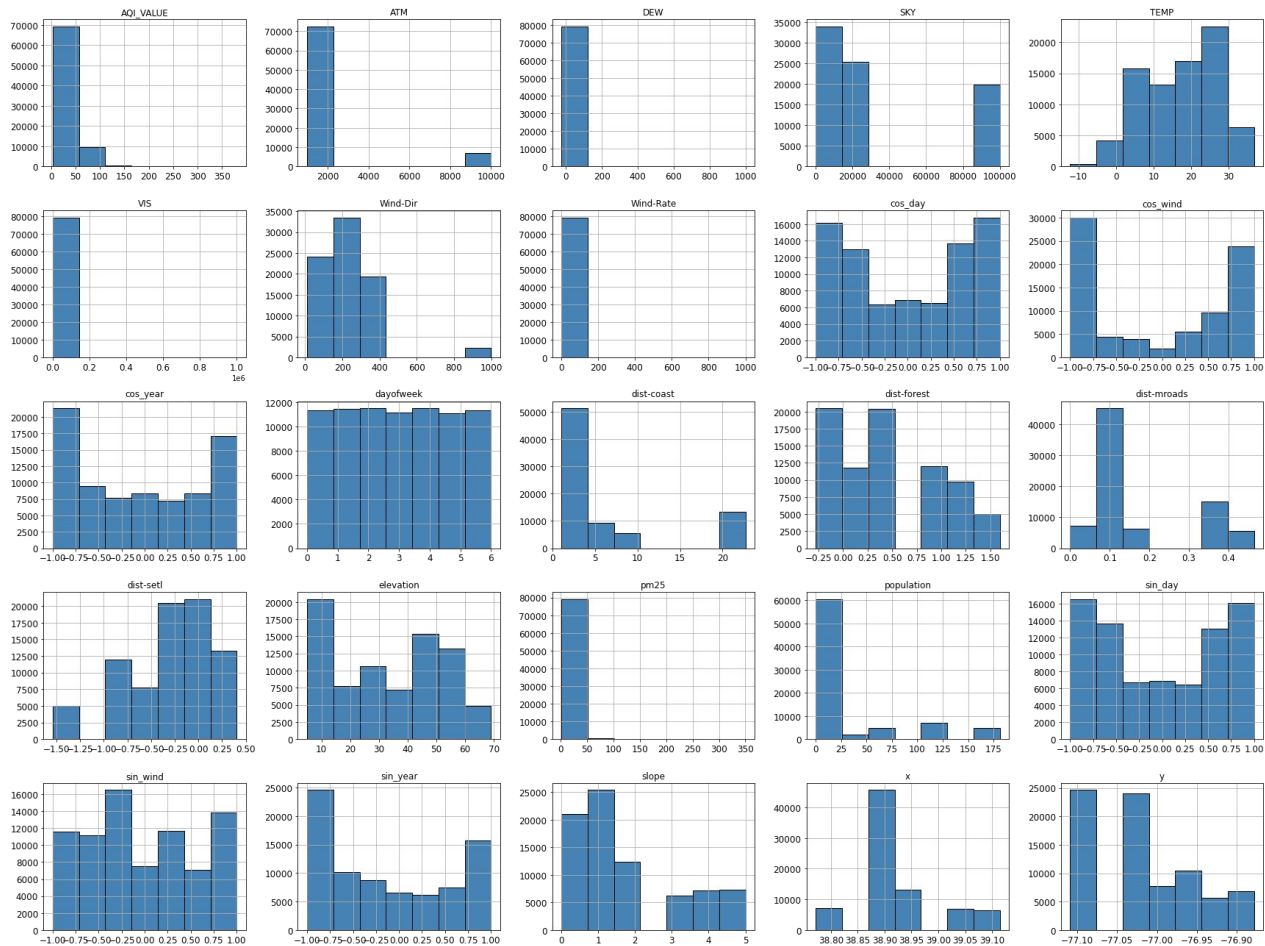


Univariate Analysis

- Goal
 - understanding of individual features within themselves
 - shape, center, and spread of individual features
- Tools
 - boxcharts, histograms, descriptive statistics
- Findings
 - Original Shape: (105398, 30)
 - Missing data
 - Outliers for pm25 and most of our weather features
- Outcomes
 - removed instances with Nan (-7724)
 - IQR to remove weather outliers (+30000)

Outlier Detection

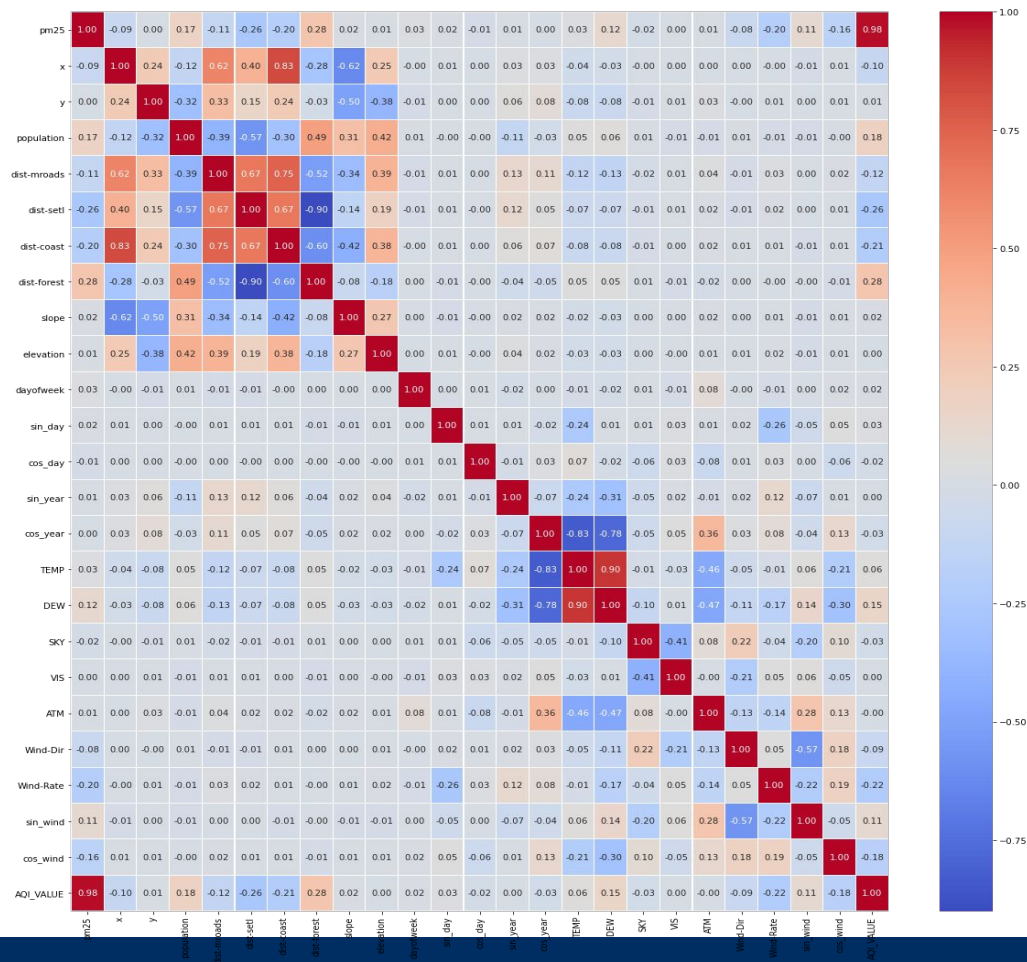




Bivariate Analysis


- Goal
 - understanding of the relationship between the features in our dataset
 - Pearson correlation
- Tools
 - correlation
 - correlation heatmaps
 - scatter plots
- Outcomes
 - eliminated any linear models for machine learning
 - helped eliminate many non-essential features (e.g. VIS, SKY)

geographic variables




Nothing is strongly correlated with our target variables (AQI and pm 2.5)


pm2.5 vs




Visibility




Distance from settlement




Pressure



Temperature



Windspeed



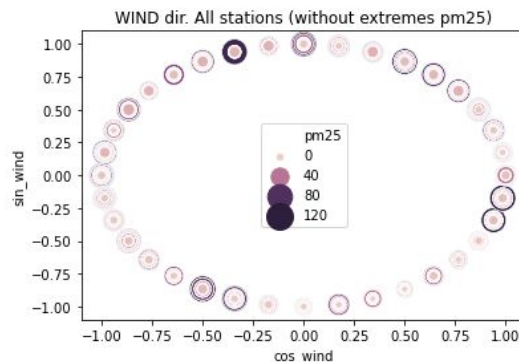
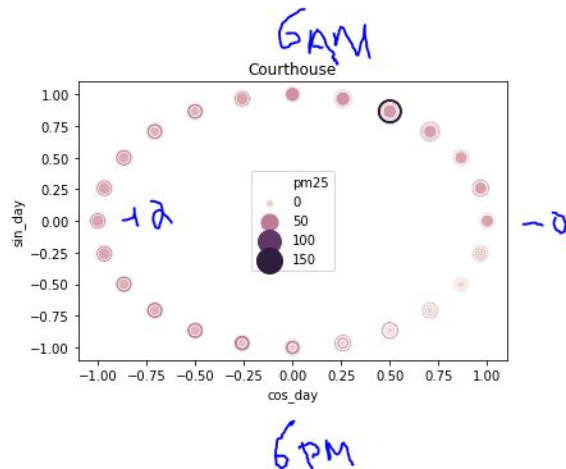
Humidity

Feature Engineering



Circular features

- datetime was splitted into 3 features :
 - sin/cos cyclic position in the year
 - sin/cos cyclic position in the day
 - category 1 to 7 for day number in the week
- Wind direction was also transformed into cos/sin cyclic




Feature Engineering

- Transformation of PM25 to EPA AQI
- Binary classification was chosen : Will the air quality be good or not.

U.S. EPA PM_{2.5} AQI

AQI Category	AQI Value	24-hr Average PM _{2.5} Concentration (µg/m ³)
Good	0 - 50	0 - 15.4
Moderate	51 - 100	15.5 - 40.4
USG	101 - 150	40.5 - 65.4
Unhealthy	151 - 200	65.5 - 150.4
Very Unhealthy	201 - 300	150.5 - 250.4
Hazardous	301 - 500	250.5 - 500.4



Model Selection and Machine Learning



Regression models



Machine Learning


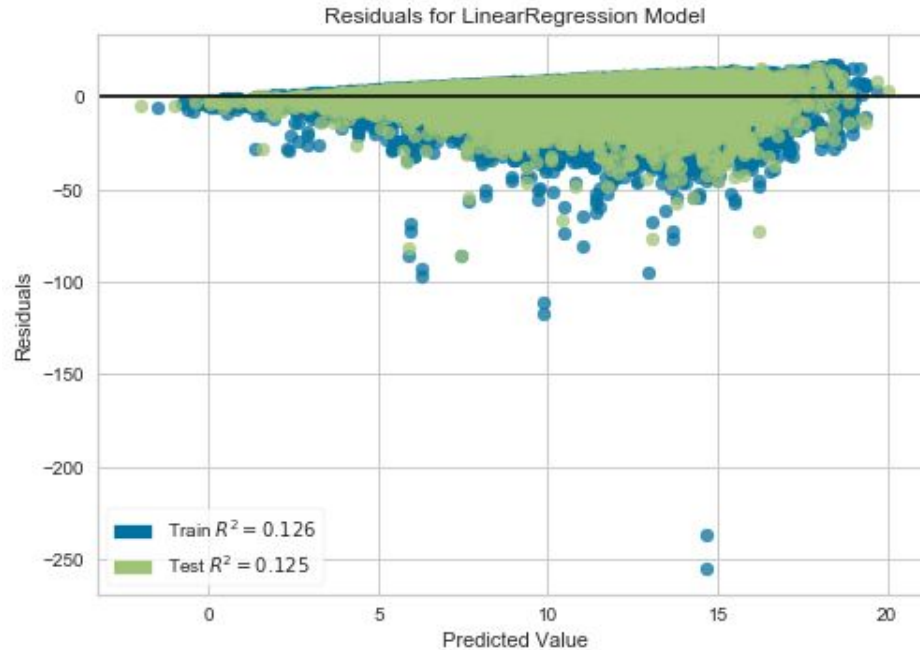


Training and Testing models


```
54]: # dictionary mapping the model classes
model_dict = {
    "DecisionTreeRegressor" : DecisionTreeRegressor,
    "MLPRegressor" : MLPRegressor,
    "KNeighborsRegressor" : KNeighborsRegressor,
    "BayesianRidge" : BayesianRidge,
    "LinearRegression" : LinearRegression
}
#list of models to be used
models = ['DecisionTreeRegressor', 'MLPRegressor', 'KNeighborsRegressor', 'BayesianRidge', 'LinearRegression']

# storing the columns to be used for training the models
# 1 -> no filter
# 2 -> using columns with little corelation
# 3 -> using columns with more corelation
process_cols = [
    ['type', 'sensor', 'x', 'y',
     'population', 'dist-mroads', 'dist-setl', 'dist-coast', 'dist-forest',
     'slope', 'elevation', 'dayofweek', 'sin_day', 'cos_day', 'sin_year',
     'cos_year', 'TEMP', 'WIND', 'DEW', 'SKY', 'VIS', 'ATM']
]
```


Regression Model Analysis



Decision Tree Regressor was overfit to the test sample



Residuals for KNeighbors Regressors Model



	model	group1	group2	group3	group4	group5	rsquare Var	rsquare mean
0	<class 'sklearn.linear_model._base.LinearRegre...	-0.702871	0.094141	-0.066153	-0.166127	0.043475	0.081909	-0.159507
1	<class 'sklearn.neighbors._regression.KNeighbo...	-0.227923	0.547875	0.301412	0.238347	0.152194	0.063629	0.202381
2	<class 'sklearn.linear_model._bayes.BayesianRi...	-0.699237	0.094671	-0.065771	-0.165153	0.043324	0.081174	-0.158433
3	<class 'sklearn.ensemble._gb.GradientBoostingR...	-0.026633	0.315593	0.218278	0.370938	0.084876	0.021452	0.192611
4	<class 'sklearn.neural_network._multilayer_per...	-18.503900	0.312136	0.141989	-1.167200	-0.068001	53.899367	-3.856995
5	<class 'sklearn.tree._classes.DecisionTreeRegr...	-0.996818	-0.169596	-0.679429	-0.223873	-0.099959	0.120529	-0.433935

Conclusion

- Best regression model among seems to be KNeighborsRegressor
- Mean of RSquare 0.2 with a variance of 0.06

We concluded that regressions were maybe too complex and because our data product can support a more simple binary classification, we decided to invest more time into classifiers

Classification models



Machine Learning



-
- Suspiciously too good results
- Overfitting or “ground station memory”

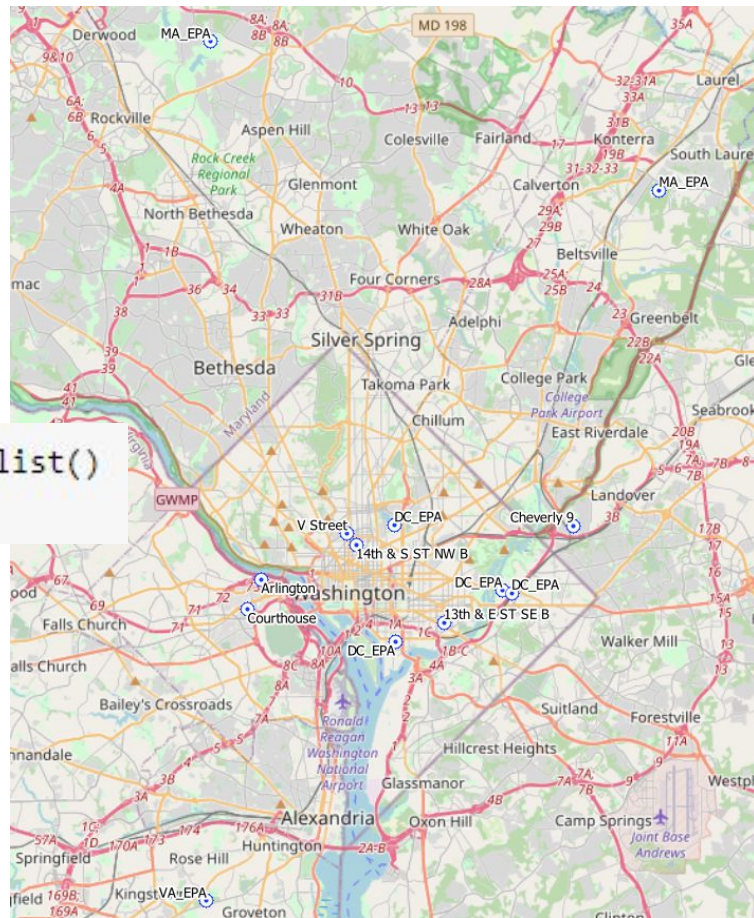
```
models = [  
    SVC(gamma='auto'),  
    # NuSVC(gamma='auto'),  
    LinearSVC(),  
    #SGDClassifier(max_iter=100, tol=1e-3),  
    KNeighborsClassifier(),  
    LogisticRegression(solver='lbfgs'), #LogisticRegressionCV(cv=3),  
    BaggingClassifier(),  
    ExtraTreesClassifier(n_estimators=100),  
    RandomForestClassifier(n_estimators=100),  
    DecisionTreeClassifier()  
]  
  
for model in models:  
    score_model(X, y, model)
```

```
SVC: 0.275444345149439  
LinearSVC: 0.0007155635062611806  
KNeighborsClassifier: 0.6596370143149284  
LogisticRegression: 0.01696885169688517  
BaggingClassifier: 0.9802424242424244  
ExtraTreesClassifier: 1.0  
RandomForestClassifier: 1.0  
DecisionTreeClassifier: 1.0
```

Machine Learning

- The test train split strategy was adjusted to group by station

```
groups = gs["station_id"].astype('category').cat.codes.tolist()
gkf = GroupKFold(n_splits=13)
```



Machine Learning



- We focused on the recall of predicting bad weather.

	model	group1	group2	group3	group4	group5	F1 NotG var	F1 NotG mean	Recall NotG var	Recall NotG mean
0	BaggingClassifier	0.280231	0.302059	0.459842	0.327219	0.665599	0.020623	0.406990	0.012572	0.477286
1	ExtraTreesClassifier	0.279841	0.340892	0.404481	0.390972	0.339450	0.001951	0.351127	0.035298	0.448685
2	RandomForestClassifier	0.297099	0.418886	0.448581	0.356493	0.356469	0.002819	0.375506	0.013719	0.395908
3	DecisionTreeClassifier	0.256872	0.259951	0.468585	0.344564	0.665886	0.023714	0.399171	0.012747	0.504688

Decision tree was actually the best model so far.

Results
Reflections,
Next Steps
Conclusion






Results

We reached 0.50 recall of predicting bad air quality so far using DecisionTreeClassifier.

We didn't find substantial better results after hyper parameter tuning.

Our first conclusions are there is potential to achieve a better modelization and we need more data.



First data product preview

One minute app

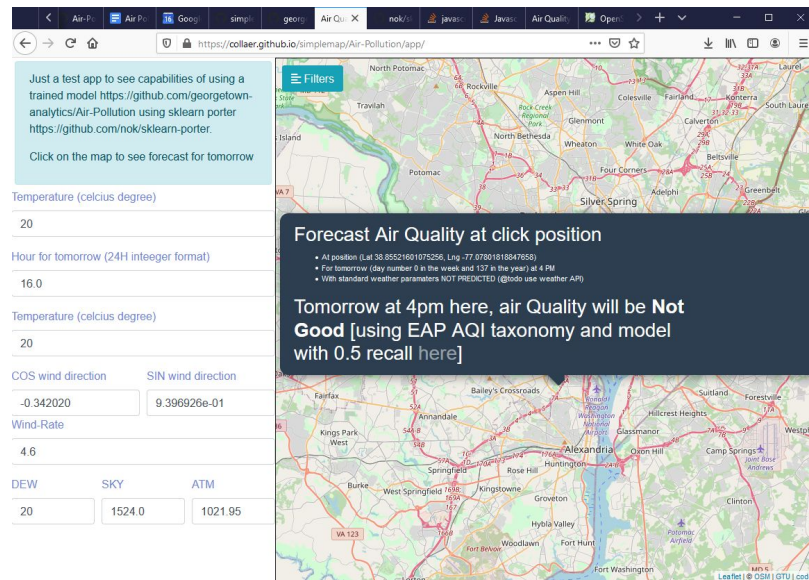
- Using <https://github.com/nok/sklearn-porter> to export model
 - Javascript file of 21000 lines
- Connecting the model.js with simple html/js app, using open source license libraries:
 - leaflet and leaflet plugins
 - bootstrap
 - jquery
- And CC open data www.openstreetmap.org
- Directly published inside Github using gh-pages branch
- gh-pages option is not activated for our repo
 - URL here: <https://georgetown-analytics.github.io/Air-Pollution/app/test>
 - Image here: <https://collaer.github.io/simplemap/Air-Pollution/app/>

```
from sklearn_porter import Porter
# Export:
porter = Porter(myTree, language='js')
output = porter.export(embed_data=True)
print(output)
```



First data product preview

```
model.js      map.js
1 | var DecisionTreeClassifier = function() {
2 |
3 |   var findMax = function(nums) {
4 |     var index = 0;
5 |     for (var i = 0; i < nums.length; i++) {
6 |       index = nums[i] > nums[index] ? i : index;
7 |     }
8 |     return index;
9 |   };
10 |
11 |   this.predict = function(features) {=};
1174 |
1175 | };
1176 |
1177 | if (typeof process !== 'undefined' && typeof process.argv !== 'undefined') {
1178 |   if (process.argv.length - 2 === 14) {
1179 |
1180 |     // Features:
1181 |     var features = process.argv.slice(2);
1182 |
1183 |     // Prediction:
1184 |     var clf = new DecisionTreeClassifier();
1185 |     var prediction = clf.predict(features);
1186 |     console.log(prediction);
1187 |
1188 |   }
1189 | }
1190 |
```



Try, Try Again

A tale of pivoting...


- Test train split data strategy, sometimes randomness is not your ally
- Data science pipeline is a loop, test, retry, removes, ...
- And of course, python stack scikit learn, yellowbrick, pandas, seaborn, etc. is awesome !



Next Steps

We believe we can reveal more information about the Air Pollution by the following improvements:
More data!

Sometimes you find something you were not looking for...
Need to dig deeper but our study suggest something about ground stations efficient management: After a period (maybe lower than a year) you can move your station to another place, because we may build a model to predict that station future values...



Thank you for
your time!



GEORGETOWN UNIVERSITY

Feedback, reactions, questions